

Adaptive Label Smoothing

```
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds
import matplotlib.pyplot as plt
import os
from pathlib import Path
from xml.etree import ElementTree
import itertools
import gc

# Use masking
# # # SET TO FALSE # # #
# # # INCORRECT USAGE # # #
MASKING = False

# Use Adaptive Label Smoothing
# 0 = Hard Label
# 1 = Full ALS
ALS = False
BETA = 0

# Train or Test
train = False
test = True

# From:
# https://www.tensorflow.org/guide/gpu#limiting\_gpu\_memory\_growth
gpus = tf.config.experimental.list_physical_devices('GPU')
if gpus:
    # Restrict TensorFlow to only allocate 1GB of memory on the first GPU
    try:
        tf.config.experimental.set_virtual_device_configuration(
            gpus[0],
            [tf.config.experimental.VirtualDeviceConfiguration(memory_limit=4096)]
        )
        logical_gpus = tf.config.experimental.list_logical_devices('GPU')
        print(len(gpus), "Physical GPUs,", len(logical_gpus), "Logical GPUs")
    except RuntimeError as e:
        # Virtual devices must be set before GPUs have been initialized
```

```

        print(e)

annotations_dir_path = os.path.join( os.getcwd(), "annotations" )
xmls_dir_path = os.path.join( annotations_dir_path, "xmls" )

unmarshallers = {

    "xmin": lambda x: int(x.text),
    "ymin": lambda x: int(x.text),
    "xmax": lambda x: int(x.text),
    "ymax": lambda x: int(x.text),
    "width": lambda x: int(x.text),
    "height": lambda x: int(x.text)

}

# Ignore the following training examples
# Missing bounding boxes
missingXMLs = {

    'samoyed_10.xml',
    'saint_bernard_15.xml',
    'Abyssinian_104.xml',
    'Ragdoll_199.xml',
    'Bengal_175.xml',
    'Bengal_111.xml',
    'Bengal_105.xml'      # Has multiple bounding boxes

}

# Constants
NUM_EPOCHS = 200
BATCH_SIZE = 16
LEARN_RATE = 0.001
IMG_SIZE = 100

# Training
NUM_CLASSES = 37
OPTIMIZER = tf.keras.optimizers.SGD( learning_rate = LEARN_RATE, momentum = 0.9 )
REGULARIZER = tf.keras.regularizers.l2( 0.01 )
LOSS = tf.keras.losses.CategoricalCrossentropy( from_logits = False )
METRIC = tf.keras.metrics.CategoricalAccuracy()

checkpoint_path = os.getcwd() + "training_1/cp-{epoch:04d}.ckpt"

```

```

checkpoint_dir = os.path.dirname(checkpoint_path)

# Create a callback that saves the model's weights from:
# https://www.tensorflow.org/tutorials/keras/save_and_load
cp_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_path,
    verbose=1,
    save_weights_only=True,
    period=20) # Saves every 10 epochs

# Decreases learning rate at specific epochs
def lrdecay(epoch):

    lr = LEARN_RATE
    if epoch > 150:
        lr *= 1e-3
    elif epoch > 100:
        lr *= 1e-2
    elif epoch > 50:
        lr *= 1e-1

    return lr

# Display some plots
def testLoad( data, info ):

    plt.figure( figsize=(10,10) )
    i=0
    for image, label in data:

        if i == 25:
            break
        plt.subplot( 5, 5, i+1 )
        plt.xticks([])
        plt.yticks([])
        plt.grid( False )
        plt.imshow( image )
        label = info.features["label"].int2str(label)
        plt.xlabel( label )
        i += 1

    plt.show()

```

```

# Architecture based on Full Pre-Activation from:
# https://arxiv.org/pdf/1603.05027.pdf
# Implementation inspired by:
# https://www.tensorflow.org/tutorials/customization/custom_layers
class identityBlock( tf.keras.Model ):

    def __init__( self, filters ):

        super( identityBlock, self ).__init__( name = '' )
        f1, f2 = filters
        k = 3    # Kernel size

        self.conv2a = tf.keras.layers.Conv2D( f1, kernel_size = (1, 1), strides =
(1, 1), padding = 'valid', kernel_regularizer = REGULARIZER )
        self.bn2a = tf.keras.layers.BatchNormalization()

        self.conv2b = tf.keras.layers.Conv2D( f1, kernel_size = (k, k), strides =
(1, 1), padding = 'same', kernel_regularizer = REGULARIZER )
        self.bn2b = tf.keras.layers.BatchNormalization()

        self.conv2c = tf.keras.layers.Conv2D( f2, kernel_size = (1, 1), strides =
(1, 1), padding = 'valid', kernel_regularizer = REGULARIZER )
        self.bn2c = tf.keras.layers.BatchNormalization()

    def call( self, inputTensor, training = False ):
        x = inputTensor

        # Block 1
        x = self.bn2a( x, training = training )
        x = tf.nn.leaky_relu( x )
        x = self.conv2a( x )

        # Block 2
        x = self.bn2b( x, training = training )
        x = tf.nn.leaky_relu( x )
        x = self.conv2b( x )

        # Block 3
        x = self.bn2c( x, training = training )
        x = tf.nn.relu( x )
        x = self.conv2c( x )

        # Output
        x += inputTensor

```

```

        return x

class convBlock( tf.keras.Model ):

    def __init__( self, filters, s ):

        super( convBlock, self ).__init__( name = '' )
        f1, f2 = filters
        k = 3    # Kernel size

        self.conv2a = tf.keras.layers.Conv2D( f1, kernel_size = (1, 1), strides =
(s, s), padding = 'valid', kernel_regularizer = REGULARIZER )
        self.bn2a = tf.keras.layers.BatchNormalization()

        self.conv2b = tf.keras.layers.Conv2D( f1, kernel_size = (k, k), strides =
(1, 1), padding = 'same', kernel_regularizer = REGULARIZER )
        self.bn2b = tf.keras.layers.BatchNormalization()

        self.conv2c = tf.keras.layers.Conv2D( f2, kernel_size = (1, 1), strides =
(1, 1), padding = 'valid', kernel_regularizer = REGULARIZER )
        self.bn2c = tf.keras.layers.BatchNormalization()

        self.conv2Shortcut = tf.keras.layers.Conv2D( f2, kernel_size = (1, 1), st
rides = (s, s), padding = 'valid', kernel_regularizer = REGULARIZER )
        self.bn2Shortcut = tf.keras.layers.BatchNormalization()

    def call( self, inputTensor, training = False ):

        x = inputTensor
        xShort = inputTensor

        # Block 1
        x = self.conv2a( x )

        # Block 2
        x = self.bn2b( x, training = training )
        x = tf.nn.leaky_relu( x )
        x = self.conv2b( x )

        # Block 3
        x = self.bn2c( x, training = training )
        x = tf.nn.relu( x )

```

```

x = self.conv2c( x )

# Shortcut
xShort = self.bn2Shortcut( xShort, training = training )
xShort = tf.nn.relu( xShort )
xShort = self.conv2Shortcut( xShort )

# Output
x += xShort

return x

# Module containing Image Classification model
# Architecture based on:
# https://towardsdatascience.com/understand-and-implement-resnet-50-with-
# tensorflow-2-0-1190b9b52691
class imgClassMod( tf.Module ):

    def __init__( self ):

        self.model = tf.keras.models.Sequential()

        # Stem Layer
        self.model.add( tf.keras.layers.ZeroPadding2D( (3, 3) ) )
        self.model.add( tf.keras.layers.Conv2D( 64, (7, 7), strides = (2, 2) ) )
        self.model.add( tf.keras.layers.BatchNormalization() )
        self.model.add( tf.keras.layers.ReLU() )
        self.model.add( tf.keras.layers.MaxPooling2D( (3, 3), strides = (2, 2) ) )
    )

    # Hidden Layers
    # Stage 1
    self.model.add( convBlock( filters = [ 64, 256 ], s = 1 ) )
    for _ in range(2):
        self.model.add( identityBlock( filters = [ 64, 256 ] ) )

    # Stage 2
    self.model.add( convBlock( filters = [ 128, 512 ], s = 2 ) )
    for _ in range(3):
        self.model.add( identityBlock( filters = [ 128, 512 ] ) )

    # Stage 3
    self.model.add( convBlock( filters = [ 256, 1024 ], s = 2 ) )
    for _ in range(5):

```

```

        self.model.add( identityBlock( filters = [ 256, 1024 ] ) )

    # Stage 4
    self.model.add( convBlock( filters = [ 512, 2048 ], s = 2 ) )
    for _ in range(2):
        self.model.add( identityBlock( filters = [ 512, 2048 ] ) )

    # Pooling
    self.model.add( tf.keras.layers.AveragePooling2D( (2, 2), padding = 'same
' ) )

    # Output
    self.model.add( tf.keras.layers.Flatten() )
    self.model.add( tf.keras.layers.Dense( NUM_CLASSES, activation = 'softmax
' ) )

    self.lrdecay = tf.keras.callbacks.LearningRateScheduler(lrdecay)    # Lea
rning rate decay
    self.model.compile( loss = LOSS,
                        optimizer = OPTIMIZER, metrics = METRIC )

    def train( self, dataGenTrain, trainImg, trainLabel, validImg, validLabel ):

        trainSteps = trainImg.shape[0] / BATCH_SIZE
        validSteps = validImg.shape[0] / BATCH_SIZE

        self.history = self.model.fit( dataGenTrain.flow( trainImg, trainLabel, b
atch_size = BATCH_SIZE ), epochs = NUM_EPOCHS,
                                       steps_per_epoch = trainSteps, validation_steps = validSteps,
                                       validation_data = ( validImg, validLabel ), callbacks = [ sel
f.lrdecay, cp_callback ] )

    def test( self, testImg, testLabel ):

        self.model.evaluate( testImg, testLabel )
        return self.model.predict( testImg )

    def load( self, pathName ):

        self.model.load_weights( pathName )

# Data Augmentation inspired by:

```

```

# https://www.tensorflow.org/tutorials/images/data_augmentation?hl=sv
def resize_and_rescale( image ):

    image = tf.cast( image, tf.float32 )
    image = tf.image.resize( image, [IMG_SIZE, IMG_SIZE] )
    image = ( image / 255.0 )

    return image

def soften( label, mask ):

    # Get sum of 1s for each bounding box == area of each bounding box
    # Divide by area of image
    alpha = tf.reduce_sum( mask, [1,2] ) / ( IMG_SIZE**2 )
    assert all( x <= 1 for x in alpha )
    return ( label * ( 1-alpha ) + ( 1-label )*alpha / ( NUM_CLASSES-
1 ) ) * BETA + ( ( 1-BETA ) * label )

def augment( image, label, mask ):

    numEx = tf.shape( image )[0]
    image = resize_and_rescale( image )

    # Add 6 pixels of padding
    image = tf.image.resize_with_crop_or_pad( image, IMG_SIZE + 6, IMG_SIZE + 6 )

    mask = tf.image.resize_with_crop_or_pad( mask, IMG_SIZE + 6, IMG_SIZE + 6 )

    # Random crop back to the original size
    seedRand = np.random.randint( 10000 )
    image = tf.image.random_crop( image, size=[ numEx, IMG_SIZE, IMG_SIZE, 3 ], s
eed=seedRand )
    mask = tf.image.random_crop( mask, size=[ numEx, IMG_SIZE, IMG_SIZE, 1 ], see
d=seedRand )

    # Adaptive Label Smoothing
    if ALS:
        label = soften( label, mask )

    return image, label, mask

# Underconfidence and Overconfidence based on:

```



```

# http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.720.9822
def uncertainty( label ):

    # Ensure not taking log of 0
    assert( np.min( labelProbs > 0 ) )
    temp = label * tf.math.log( label ) / tf.math.log( tf.convert_to_tensor( NUM_
CLASSES, dtype=tf.float32 ) ) # Convert to log base NUM_CLASSES
    return -tf.reduce_sum( temp, axis=-1 ) # Sum for each example

def calcConf( labelPreds, labelTruthSparse, ind, labelProbs, labelTruth ):

    numEx = tf.shape( labelProbs )[0]

    K_c = tf.cast( tf.reduce_sum( ind ), tf.int32 ) # Gets number of corr
ect predictions
    K_f = numEx - K_c

    h = uncertainty( labelProbs )
    confCorrect = ind * h
    confIncorrect = ( 1-ind ) * ( 1-h )
    underConf = tf.cast( 1 / K_c, tf.float32 ) * tf.reduce_sum( confCorrect )
    overConf = tf.cast( 1 / K_f, tf.float32 ) * tf.reduce_sum( confIncorrect )

    return underConf, overConf

# ECE and MCE based on:
# https://arxiv.org/pdf/1706.04599.pdf and
# http://people.cs.pitt.edu/~milos/research/AAAI_Calibration.pdf
def calcCE( labelPreds, labelTruthSparse, ind, labelProbs, labelTruth, binSize ):

    numEx = tf.shape( labelProbs )[0]

    # Split into bins
    unevenSize = numEx.numpy() % binSize # Get the leftover bin
    labelProbs, unevenProbs = tf.split( labelProbs, [ numEx-
unevenSize, unevenSize ] )
    ind, unevenInd = tf.split( ind, [ numEx-unevenSize, unevenSize ] )
    splitProbs = tf.split( labelProbs, binSize )
    splitAcc = tf.split( ind, binSize )

    # For easier computation (list -> tensor)
    splitProbs = tf.transpose( splitProbs, [1,0,2] )
    splitAcc = tf.transpose( splitAcc )

```

```

    # Uneven calculations
    unevenAcc = tf.reduce_sum( unevenInd, axis=-
1 ) / unevenSize          # Gets number of correct predictions
    unevenConf = tf.reduce_sum( tf.math.reduce_max( unevenProbs, axis=-
1 ), axis=-1 ) # Gets max probability
    unevenConf = unevenConf / unevenSize

    # Bin calculations
    binAcc = tf.reduce_sum( splitAcc, axis=-
1 ) / binSize          # Gets number of correct predictions
    binConf = tf.reduce_sum( tf.math.reduce_max( splitProbs, axis=-1 ), axis=-
1 ) # Gets max probability
    binConf = binConf / binSize

    unevenDiff = abs( unevenAcc - unevenConf )
    binDiff = abs( binAcc - binConf )

    ECE = ( tf.reduce_sum( unevenDiff*unevenSize ) + tf.reduce_sum( binDiff*binSi
ze, axis=-1 ) ) / tf.cast( numEx, tf.float32 )
    MCE = max( tf.reduce_max( unevenDiff ), tf.reduce_max( binDiff ) )

    return ECE, MCE

def calcStats( labelProbs, labelTruth ):

    labelPreds = tf.math.argmax( labelProbs, axis=-
1 )          # Gets sparse prediction for each example
    labelTruthSparse = tf.math.argmax( labelTruth, axis=-
1 ) # Gets sparse label for each example
    ind = ( labelPreds.numpy() == labelTruthSparse.numpy() ) # Get number of c
orrect labels
    ind = tf.cast( ind, tf.float32 )

    underConf, overConf = calcConf( labelPreds, labelTruthSparse, ind, labelProbs
, labelTruth )
    ECE15, MCE = calcCE( labelPreds, labelTruthSparse, ind, labelProbs, labelTrut
h, 15 )
    ECE100, _ = calcCE( labelPreds, labelTruthSparse, ind, labelProbs, labelTruth
, 100 )

    # Display results
    print( "O.Conf: ", overConf )
    print( "U.Conf: ", underConf )

```

```

print( "ECE15: ", ECE15 )
print( "ECE100: ", ECE100 )
print( "MCE: ", MCE )

return ECE15, ECE100, MCE

if __name__ == "__main__":

    # Load data
    (dsTrain, dsValid), info = tfds.load( 'oxford_iiit_pet', split=[ 'train', 'test' ], with_info=True )

    # Test images
    #testLoad( dsTrain, info )
    #testLoad( dsValid, info )

    # Load data into a dictionary
    trainDict = { 'image': [], 'label': [], 'mask': [] }
    validDict = { 'image': [], 'label': [] }

    for example in dsTrain:

        image_name = example["file_name"]
        image = example["image"]
        label = example["label"]

        # Get name of file
        name = tf.get_static_value( image_name ).decode( 'utf-8' )
        name = os.path.splitext( name )[0]
        xml_name = name + ".xml"

        if( xml_name in missingXMLs ):
            continue

        # Add bounding boxes
        xml_path = os.path.join( xmls_dir_path, xml_name )
        sizeInfo = []
        for _, elem in ElementTree.iterparse( xml_path ):

            unmarshal = unmarshallers.get(elem.tag)
            if unmarshal:
                data = unmarshal(elem)
                elem.clear()
                elem.text = data

```

```

        sizeInfo.append( elem.text )

    assert ( len(sizeInfo) == 6 )    # Ensure there aren't multiple bounding boxes

    W, H, w1, h1, w2, h2 = sizeInfo

    # Create mask
    mask = np.zeros( ( H,W ) )
    mask[ h1:h2, w1:w2 ] = 1
    mask = np.expand_dims( mask, -1 )

    # Resize images
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
    mask = tf.image.resize(mask, [IMG_SIZE, IMG_SIZE])

    trainDict[ 'image' ].append( image )
    trainDict[ 'label' ].append( label )
    trainDict[ 'mask' ].append( mask )

for example in dsValid:

    image = example[ 'image' ]
    label = example[ 'label' ]
    image = resize_and_rescale( image )
    validDict[ 'image' ].append( image )
    validDict[ 'label' ].append( label )

# Clear memory
dsTrain = None
dsValid = None
info = None
gc.collect()

print( "Number of training examples:", len( trainDict[ 'image' ] ) )
assert all( x.shape == (IMG_SIZE, IMG_SIZE, 3) for x in trainDict[ 'image' ]
)

# Combine list of arrays into a single array
trainImg = tf.stack( trainDict[ 'image' ], axis=0 )
trainLabel = np.hstack( trainDict[ 'label' ] )
validImg = tf.stack( validDict[ 'image' ], axis=0 )
validLabel = np.hstack( validDict[ 'label' ] )

mask = trainDict[ 'mask' ]

```

```

# Clear memory
trainDict = None
validDict = None
gc.collect()

# One-Hot Encode for ALS
trainLabel = tf.one_hot( trainLabel, NUM_CLASSES )
validLabel = tf.one_hot( validLabel, NUM_CLASSES )

# Apply random cropping to training data
trainImg, trainLabel, mask = augment( trainImg, trainLabel, mask )

if MASKING:

    # Compute pixel-mean mask
    avgVal = tf.reduce_mean( trainImg, [ 1, 2 ] )
    avgVal = tf.reshape( avgVal, [ -1, 1, 1, 3 ] )
    maskAdd = (1-mask) * avgVal.numpy()

    # Apply masking
    trainImg = trainImg * mask
    trainImg = trainImg + maskAdd

mask = None
maskAdd = None
gc.collect()

# Test images
for i in range(25):

    plt.subplot( 5, 5, i+1 )
    plt.imshow( trainImg[i] )

plt.show()

# Convert to tensor
trainImg = tf.convert_to_tensor( trainImg, dtype=tf.float32 )
trainLabel = tf.convert_to_tensor( trainLabel )
validImg = tf.convert_to_tensor( validImg, dtype=tf.float32 )
validLabel = tf.convert_to_tensor( validLabel )

# Image Data Generator
dataGenTrain = tf.keras.preprocessing.image.ImageDataGenerator( rotation_range=15, width_shift_range=0.2, height_shift_range=0.2, horizontal_flip=True )
dataGenTrain.fit( trainImg )

```

```

# Ensure all values are normalized
assert all( np.max(x) <= 1 for x in trainImg )
assert all( np.max(x) <= 1 for x in validImg )

# Model
model = imgClassMod()

if train:

    model.train( dataGenTrain, trainImg, trainLabel, validImg, validLabel )

if test:

    pathName = os.getcwd() + "\\Checkpoints\\"
    pathName += "Baseline"
    pathName += "\\cp-0200.ckpt"
    print(pathName)
    model.load( pathName )

    labelProbs = tf.convert_to_tensor( model.test( validImg, validLabel ) )

# Calculate metrics
calcStats( labelProbs, validLabel )

```

I used the Oxford IIIT Pets dataset. This dataset consists of around 200 images per class for 37 classes. However, the training and test sets were created with a 50-50 split. I initially wanted to recombine them and create my own 80-20 or 90-10 split, but I found that the images in the test set did not have bounding boxes, so I could not use them for training. This meant that for training, I have around 100 images per class for 37 classes, which is a very small amount. I suspected this would easily overfit, so I attempted to mitigate this by using TensorFlow's ImageDataGenerator to augment the training examples, creating new examples.

I decided to exclude some examples for two main reasons. The first reason is that some of the images were unusable; They were either saved under the incorrect format (.jpg instead of

.png or .gif), or they were corrupted. The second reason is that some of the training images did not have proper bounding boxes. Some of them did not have an .xml file, meaning there was no associated bounding box for the image. One of the images had two bounding boxes, and the authors mentioned a strategy (masking) for dealing with images with multiple bounding boxes. However, as there was only one image with more than one bounding box, I decided it would not significantly change the results if it was not included. I excluded this image for consistency and convenience.

I decided not to use two validation sets like the authors did. This is because the V2 validation set that the authors used was a set of “challenging” images compared to the V1 validation set. However, the Oxford IIT Pets dataset does not have a comparable split. Randomly splitting the validation set into two independent sets would have been pointless, as they would both have similar results.

I initially misunderstood what the authors meant by ‘masking’. I initially understood it to mean replacing every pixel outside of the bounding box with a constant value. However, after getting terrible results, I realized that the authors use it to deal with images with multiple bounding boxes. The authors mask every bounding box except for one, which also results in additional training examples. This is not very applicable to the dataset I used, because all of the examples except for one have a single bounding box. I replicated the masking results achieved by the authors without using masking.

I resize every image to 100 by 100 pixels. I initially resized to 150 by 150 pixels, but I decided to further reduce the size to facilitate faster training. I also add 6 pixels to each image and randomly crop it back to 100 by 100 pixels.

I use ResNet50. I train for 200 epochs using SGD with a learning rate of 0.001 decayed by 0.1 at epochs 50, 100, and 150 and a momentum of 0.9. I use a batch size of 16 and L2 regularization with a penalty term of 0.01.

At the end of training, my model achieves over 99% training accuracy, while my validation accuracy is around 47%. Not only is there significant overfitting, but the almost perfect training accuracy suggests that the model is memorizing the training data. I did not take additional measures to combat overfitting as that was not the primary goal of the paper and because the limited examples available in the dataset were likely to result in overfitting.

Method	Train N	ACC	ECE 100	ECE 15	MCE	O.conf	U.conf
Hard Label	3.668k	0.465	0.289	0.289	0.631	0.728	0.127
A. L. S.	3.668k	0.438	0.036	0.097	0.368	0.285	0.568
A. L. S. (Beta=0.75)	3.668k	0.457	0.031	0.102	0.460	0.331	0.515
A. L. S. (Beta=0.25)	3.668k	0.460	0.132	0.147	0.460	0.510	0.321
Hard Label (incorrect mask)	3.668k	0.148	N/A	N/A	N/A	N/A	N/A
A. L. S. (incorrect mask)	3.668k	0.227	N/A	N/A	N/A	N/A	N/A

Table 1: Classification and calibration results with Oxford IIIT Pets

The results in Table 1 are mostly consistent with the results achieved in the paper. The Hard Label method is equivalent to $\beta=0$. The A.L.S. method is equivalent to $\beta=1$. With a decreasing β , the overconfidence increases while the underconfidence decreases. The accuracies achieved are also consistent. The accuracies are mostly consistent, with a slight increase with a decreasing β .

MCE was calculated along with ECE15. The authors did not mention the bin size used to calculate MCE. Calculating MCE along with ECE100 resulted in a smaller value. It is possible that there is an error in the implementation of the ECE and MCE calculations. The results obtained are the opposite of the results from the paper. In the results, ECE100 is consistently smaller than ECE15, while in the paper, the opposite is true. In the results, ECE and MCE increase with a decreasing beta, while in the paper, the opposite is true.

The results obtained by masking are reported. The Hard Label (incorrect mask) method utilized zero masking, meaning the pixels outside of the bounding box were masked with zeroes. The A.L.S. (incorrect mask) method utilized mean pixel masking, meaning the pixels outside of the bounding box were masked with the average pixel value for each respective channel. The additional statistics were not computed for these incorrect methods.

A.L.S. with $\beta=0.75$ appears to be the best compromise out of all of the methods. It has a small reduction in accuracy, while significantly reducing the overconfidence of the model. This is in contrast with A.L.S. with $\beta=1$, which has a larger decrease in accuracy while barely reducing the overconfidence compared to A.L.S. with $\beta=0.75$. A.L.S. with $\beta=0.25$ achieves almost the same accuracy as A.L.S. with $\beta=0.75$, while also having a significantly higher overconfidence.

Overall, a lot of my problems were caused by the dataset I selected. However, I chose the dataset, so I had to work through the problems. It is very likely that I could have achieved better results with a better dataset. If I had more computing power, I could have resized the images to a large size. Currently, resizing to 100 by 100 pixels discards a lot of information for the larger images.