

```
import struct as st
import numpy as np
import tensorflow as tf
import kerastuner as kt
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import os

# Constants
NUM_EPOCHS = 30
VALIDATION_PERCENTAGE = 0.2
DROP_RATE = 0.2

# Image specific parameters
IMG_DIMS = ( 28, 28, )
NUM_CLASSES = 10
HIDDEN_SIZE = [ 512, 512 ] # Number of output neurons for each hidden layer

# Data processing inspired by:
# https://stackoverflow.com/questions/39969045/parsing-yann-lecuns-mnist-idx-
# file-format
def imgRead( fileName ):

    with open( fileName, 'rb') as f:

        # Go to beginning of file and start reading
        f.seek(0)
        _numMagic = st.unpack( '>I', f.read(4) )[0] # Read magic
number
        numImg, numRows, numCols = st.unpack( '>III', f.read(12) ) # Read ints

        numTotalBytes = numImg * numRows * numCols * 1 # Each pixel is 1 byte

        # Read the remaining data in the file
        img = np.asarray( st.unpack(
            '>' + 'B'*numTotalBytes, f.read( numTotalBytes )
        ) ).reshape( ( numImg, numRows, numCols ) )

    return img
```

```

def labelRead( fileName ):

    with open( fileName, 'rb' ) as f:

        # Go to beginning of file and start reading
        f.seek(0)
        _numMagic = st.unpack( '>I', f.read(4) )[0]      # Read magic number
        numImg = st.unpack( '>I', f.read(4) )[0]        # Read number of examples

        # Read the remaining data in the file
        labels = st.unpack( '>' + 'B'*numImg, f.read( numImg ) )

    return labels


def model_builder( hp ):

    model = tf.keras.models.Sequential()

    # Input Layer
    model.add( tf.keras.layers.Flatten( input_shape = IMG_DIMS ) )

    # Hidden Layers
    for numNeurons in HIDDEN_SIZE:
        model.add( tf.keras.layers.Dense( numNeurons, activation = 'relu', kernel
_regularizer = 'l2' ) )
        model.add( tf.keras.layers.Dropout( hp.Choice( 'rate', values = [ 0.3, 0.
25, 0.2 ] ), trainable = False ) )

    # Output Layer
    model.add( tf.keras.layers.Dense( NUM_CLASSES, activation = 'softmax' ) )

    # Compile model
    optimizer = tf.keras.optimizers.SGD(
        hp.Choice( 'learning_rate', values = [ 0.01, 0.005 ] ) )
    model.compile( loss = "sparse_categorical_crossentropy",
        optimizer = optimizer, metrics = "accuracy" )

    return model


def plotAccuracy( history ):

    # Plots accuracy over time
    plt.figure( figsize = (10,5) )
    plt.plot( history.history[ 'accuracy' ] )

```

```

plt.plot( history.history['val_accuracy'] )
plt.title( 'Model Accuracy' )
plt.xlabel( 'Epochs' )
plt.ylabel( 'Accuracy', rotation = 'horizontal', ha = 'right' )
plt.legend( [ 'Train', 'Valid' ], loc = 'upper left' )
plt.show()

def main():

    # Load data
    img = imgRead( 'trainImages.idx3-ubyte' )
    label = labelRead( 'trainLabels.idx1-ubyte' )
    testImg = imgRead( 'testImages.idx3-ubyte' )
    testLabel = labelRead( 'testLabels.idx1-ubyte' )

    trainImg, validImg, trainLabel, validLabel = train_test_split( img, label, te
st_size = 0.2 )

    # Normalize and convert to tensor
    trainImg = tf.convert_to_tensor( trainImg / 255, dtype=tf.float32 )
    trainLabel = tf.convert_to_tensor( trainLabel )
    validImg = tf.convert_to_tensor( validImg / 255, dtype=tf.float32 )
    validLabel = tf.convert_to_tensor( validLabel )
    testImg = tf.convert_to_tensor( testImg / 255, dtype=tf.float32 )
    testLabel = tf.convert_to_tensor( testLabel )

    # Test if images were properly loaded
    plt.figure()
    plt.imshow( trainImg[0,:,:], cmap = 'gray' )
    plt.show()

    # Test if labels were properly loaded
    print( trainLabel[0] )

    # Initialize model
    tuner = kt.Hyperband(

        model_builder,
        objective = "val_accuracy",
        max_epochs = 10,
        directory = os.path.normpath( 'D:/ ' ),
        project_name = 'tunedParams'

    )

```

```

tuner.search(

    trainImg, trainLabel, epochs = 10,
    validation_data = (validImg, validLabel)

)

best_hps = tuner.get_best_hyperparameters( num_trials = 2 )[0]

# Train and Test
print( "\nStarted Training\n" )
classifier = tuner.hypermodel.build( best_hps )
history = classifier.fit(
    trainImg, trainLabel, epochs = NUM_EPOCHS,
    validation_data = ( validImg, validLabel ) )
print( "\nFinished Training\n" )
classifier.evaluate( testImg, testLabel )

# Display accuracy
plotAccuracy( history )

main()

```

Training and Testing were performed using the default batch size of 32. The learning rate of SGD, along with the drop rate for each dropout layer was learned through the Keras Tuner.

The validation set was created with an 80/20 split of the original training set. Training set had 48,000 examples and validation set had 12,000 examples.

Hidden layers use L2 regularization with the default regularization parameter of 0.01.




Figure 1: Example plot for the training and validation accuracy of the model during training

313/313 [=====] - 0s 2ms/step - loss: 0.3282 - accuracy: 0.9598

Figure 2: Final accuracy on the test set